

Hypothesis_Testing_Cars_dataset

November 13, 2025

1 Hypothesis Testing on Cars Dataset

1.1 Awais Manzoor

2 Import Libraries

```
[8]: import pandas as pd
import numpy as np
from scipy import stats
import seaborn as sns
import matplotlib.pyplot as plt
```

3 Load Dataset

```
[7]: cars = pd.read_excel('CARS.xlsx')
cars.head(10)
```

```
[7]:      Make          Model   Type Origin DriveTrain   MSRP Invoice \
0   Acura        MDX     SUV   Asia    All  36945  33337
1   Acura       RSX Type S 2dr  Sedan   Asia  Front  23820  21761
2   Acura        TSX 4dr  Sedan   Asia  Front  26990  24647
3   Acura        TL 4dr  Sedan   Asia  Front  33195  30299
4   Acura     3.5 RL 4dr  Sedan   Asia  Front  43755  39014
5   Acura  3.5 RL w/Navigation 4dr  Sedan   Asia  Front  46100  41100
6   Acura     NSX coupe 2dr manual S  Sports   Asia   Rear  89765  79978
7   Audi        A4 1.8T 4dr  Sedan Europe  Front  25940  23508
8   Audi  A41.8T convertible 2dr  Sedan Europe  Front  35940  32506
9   Audi        A4 3.0 4dr  Sedan Europe  Front  31840  28846

      EngineSize  Cylinders  Horsepower  MPG_City  MPG_Highway  Weight \
0            3.5         6.0         265        17             23    4451
1            2.0         4.0         200        24             31    2778
2            2.4         4.0         200        22             29    3230
3            3.2         6.0         270        20             28    3575
4            3.5         6.0         225        18             24    3880
5            3.5         6.0         225        18             24    3893
6            3.2         6.0         290        17             24    3153
```

7	1.8	4.0	170	22	31	3252
8	1.8	4.0	170	23	30	3638
9	3.0	6.0	220	20	28	3462

	Wheelbase	Length
0	106	189
1	101	172
2	105	183
3	108	186
4	115	197
5	115	197
6	100	174
7	104	179
8	105	180
9	104	179

[10]: cars.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 428 entries, 0 to 427
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --    
 0   Make         428 non-null    object  
 1   Model        428 non-null    object  
 2   Type          428 non-null    object  
 3   Origin        428 non-null    object  
 4   DriveTrain   428 non-null    object  
 5   MSRP          428 non-null    int64  
 6   Invoice       428 non-null    int64  
 7   EngineSize   428 non-null    float64 
 8   Cylinders    426 non-null    float64 
 9   Horsepower   428 non-null    int64  
 10  MPG_City     428 non-null    int64  
 11  MPG_Highway 428 non-null    int64  
 12  Weight        428 non-null    int64  
 13  Wheelbase    428 non-null    int64  
 14  Length        428 non-null    int64  
dtypes: float64(2), int64(8), object(5)
memory usage: 50.3+ KB
```

- load the dataset and check columns, data types, and missing values.

4 Data Cleaning

```
[11]: cars.isnull().sum()
```

```
[11]: Make      0  
Model      0  
Type       0  
Origin     0  
DriveTrain 0  
MSRP       0  
Invoice    0  
EngineSize 0  
Cylinders  2  
Horsepower 0  
MPG_City   0  
MPG_Highway 0  
Weight     0  
Wheelbase  0  
Length     0  
dtype: int64
```

```
[13]: # Handle missing values in Cylinders  
cars['Cylinders'].fillna(cars['Cylinders'].median(), inplace=True)
```

C:\Users\hp\AppData\Local\Temp\ipykernel_17696\3393017961.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
cars['Cylinders'].fillna(cars['Cylinders'].median(), inplace=True)
```

```
[14]: print(cars.isnull().sum())
```

```
Make      0  
Model      0  
Type       0  
Origin     0  
DriveTrain 0  
MSRP       0  
Invoice    0  
EngineSize 0
```

```
Cylinders      0
Horsepower     0
MPG_City       0
MPG_Highway    0
Weight          0
Wheelbase       0
Length          0
dtype: int64
```

- We replace missing values (only 2 in Cylinders) with the median — suitable for numerical data.

5 Descriptive Statistics

```
[15]: cars.describe()
```

```
[15]:          MSRP        Invoice   EngineSize  Cylinders  Horsepower \
count    428.000000  428.000000  428.000000  428.000000  428.000000
mean    32774.855140 30014.700935   3.196729    5.808411  215.885514
std     19431.716674 17642.117750   1.108595    1.554844  71.836032
min    10280.000000  9875.000000   1.300000    3.000000  73.000000
25%    20334.250000 18866.000000   2.375000    4.000000 165.000000
50%    27635.000000 25294.500000   3.000000    6.000000 210.000000
75%    39205.000000 35710.250000   3.900000    6.000000 255.000000
max    192465.000000 173560.000000   8.300000   12.000000 500.000000

          MPG_City  MPG_Highway      Weight   Wheelbase      Length
count    428.000000  428.000000  428.000000  428.000000  428.000000
mean     20.060748   26.843458  3577.953271  108.154206  186.362150
std      5.238218   5.741201  758.983215   8.311813  14.357991
min     10.000000   12.000000  1850.000000  89.000000 143.000000
25%    17.000000   24.000000  3104.000000 103.000000 178.000000
50%    19.000000   26.000000  3474.500000 107.000000 187.000000
75%    21.250000   29.000000  3977.750000 112.000000 194.000000
max    60.000000   66.000000  7190.000000 144.000000 238.000000
```

- This gives the mean, median, min, max, std for numeric columns like MSRP, Horsepower, EngineSize, etc.
- It helps understand data distribution before testing hypotheses.

6 Normality

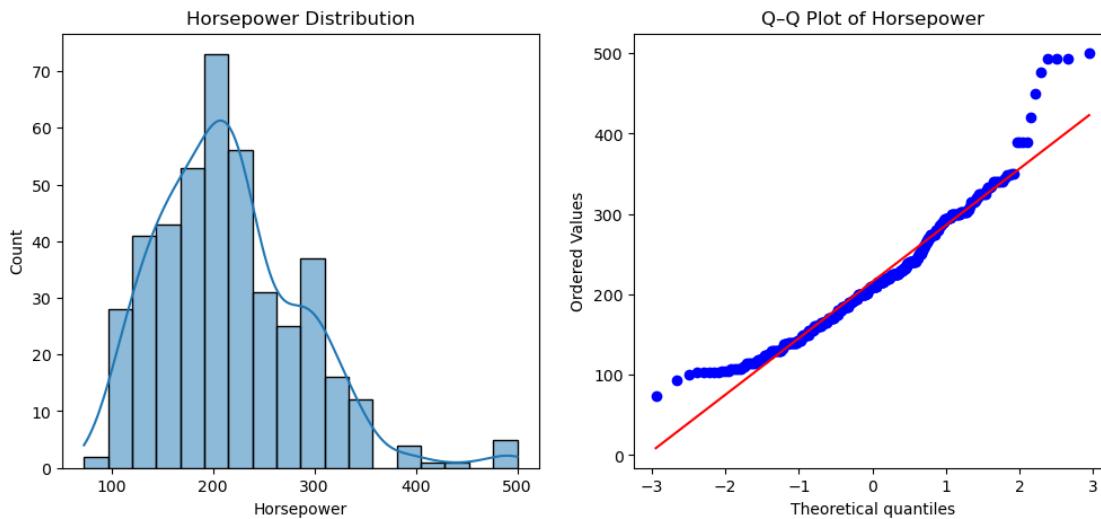
- before Parametric Tests

6.1 Histogram & Q-Q Plot

```
[17]: import scipy.stats as stats

plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
sns.histplot(cars['Horsepower'], kde=True)
plt.title("Horsepower Distribution")

plt.subplot(1,2,2)
stats.probplot(cars['Horsepower'], dist="norm", plot=plt)
plt.title("Q-Q Plot of Horsepower")
plt.show()
```



7 Parametric Tests

- when data is normal

8 Hypothesis Tests Overview

Test	Use Case	Hypothesis Example	Decision Rule
One-Sample t-test	Compare sample mean to value	$H : \text{Horsepower} = 200$	Reject / Fail based on $p < 0.05$
Independent t-test	Compare two independent groups	$H : \text{Mean MSRP (USA)} = \text{Mean MSRP (Europe)}$	Based on p-value significance

Test	Use Case	Hypothesis Example	Decision Rule
ANOVA	Compare 3+ groups	$H : \text{Mean Horsepower} (\text{Sedan} = \text{SUV} = \text{Sports})$	$p < 0.05 \rightarrow \text{at least one mean differs}$
Mann–Whitney U	Non-parametric 2-group comparison	$H : \text{Distributions equal} (\text{USA} vs \text{Europe})$	Based on p-value significance
Kruskal–Wallis	Non-parametric ANOVA (3+ groups)	$H : \text{All medians equal}$ (car types)	$p < 0.05 \rightarrow \text{at least one differs}$
Chi-Square	Relationship between categories	$H : \text{Origin and DriveTrain independent}$	$p < 0.05 \rightarrow \text{variables related}$

9 One-Sample t-test

- checking if there is a significant difference between a sample and hypothesized population means.

```
[26]: from scipy.stats import ttest_1samp
```

```
[18]: t_stat, p_val = stats.ttest_1samp(cars['Horsepower'], 200)
print("t-value:", t_stat)
print("p-value:", p_val)

if p_val < 0.05:
    print("Reject H → Mean horsepower is significantly different from 200.")
else:
    print("Fail to reject H → Mean horsepower = 200.")
```

t-value: 4.574891766789726
p-value: 6.253344784284048e-06
Reject H → Mean horsepower is significantly different from 200.

10 Independent Two-Sample t-test

- A two-sample t-test is used for comparing the significant difference between two independent groups. This test is also known as an independent samples t-test.

```
[25]: from scipy.stats import ttest_ind
```

```
[20]: usa = cars[cars['Origin'] == 'USA']['MSRP']
europe = cars[cars['Origin'] == 'Europe']['MSRP']

t_stat, p_val = stats.ttest_ind(usa, europe)
print("t-value:", t_stat)
print("p-value:", p_val)
```

```

if p_val < 0.05:
    print("Reject H → Mean MSRP differs between USA and Europe cars.")
else:
    print("Fail to reject H → No significant difference in MSRP.")

```

t-value: -8.5368542432442
p-value: 1.0436671157614213e-15
Reject H → Mean MSRP differs between USA and Europe cars.

11 ANOVA (*One-Way*)

- With the one-way ANOVA method, we compare multiple groups based on only one independent variable

[22]: `from scipy.stats import f_oneway`

```

[24]: stat, p=f_oneway(cars["Invoice"],cars["MSRP"],cars["Horsepower"])
print("p_value:",p)
print("ANOVA:",stat)
if p < 0.05:
    print("Reject H → At least one group mean is different.")
else:
    print("Fail to reject H → All group means are equal.")

```

p_value: 2.7292133264922565e-186
ANOVA: 607.5629296531739
Reject H → At least one group mean is different.

12 Non-Parametric Tests

- if data not normal

13 Mann–Whitney U Test

- Alternative to independent t-test (used when data is not normal).

```

[27]: stat, p_val = stats.mannwhitneyu(usa, europe)
print("U-statistic:", stat)
print("p-value:", p_val)
if p_val < 0.05:
    print("Reject H → Median MSRP differs between USA and Europe cars.")
else:
    print("Fail to reject H → No significant difference in median MSRP.")

```

U-statistic: 3489.5
p-value: 3.749254035965752e-18
Reject H → Median MSRP differs between USA and Europe cars.

14 Kruskal–Wallis Test

- Non-parametric version of ANOVA.

```
[29]: groups = [grp['MSRP'].values for name, grp in cars.groupby('Type')]
stat, p_val = stats.kruskal(*groups)
print("Statistic:", stat)
print("p-value:", p_val)
if p_val < 0.05:
    print("Reject H → At least one car type has a different median MSRP.")
else:
    print("Fail to reject H → Median MSRP is similar across car types.")
```

Statistic: 49.18863253683512
p-value: 2.030713990529349e-09
Reject H → At least one car type has a different median MSRP.

15 Chi-Square Test of Independence

```
[32]: from scipy.stats import chi2_contingency

[33]: contingency_table = pd.crosstab(cars['Origin'], cars['DriveTrain'])
chi2, p, dof, expected = chi2_contingency(contingency_table)
print("Chi2:", chi2)
print("p-value:", p)

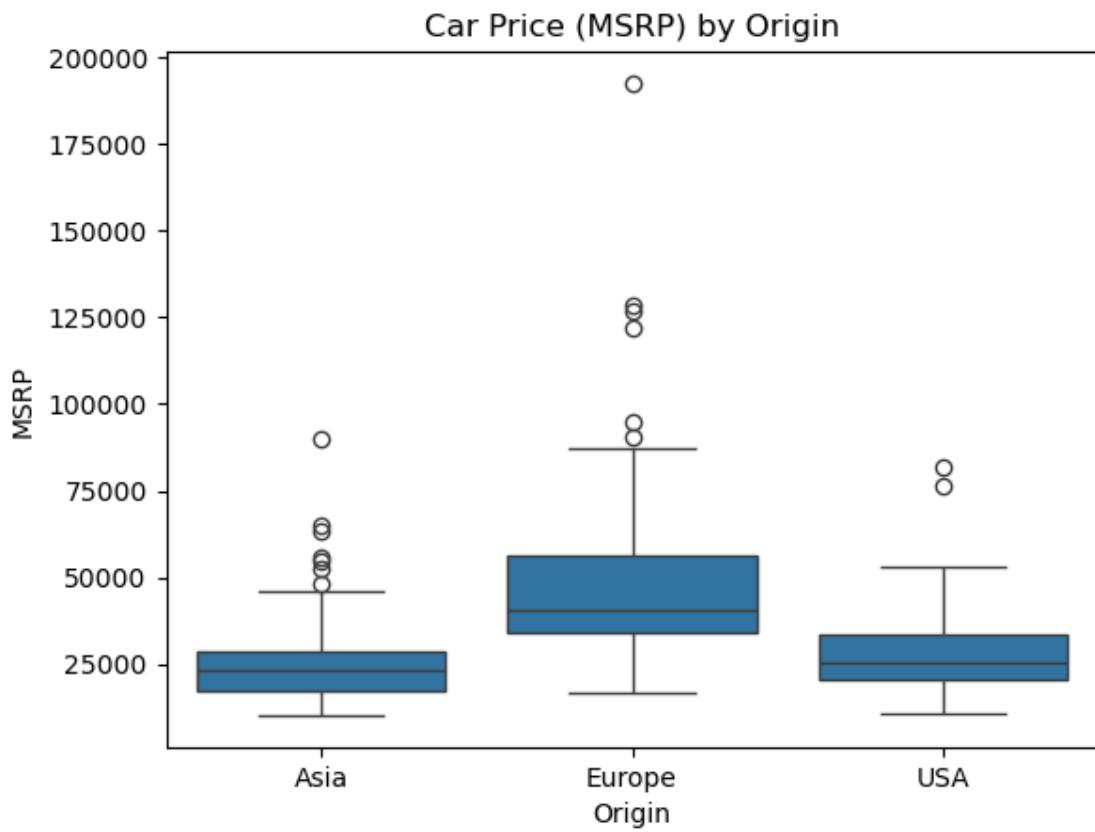
if p < 0.05:
    print("Reject H → Origin and DriveTrain are related.")
else:
    print("Fail to reject H → No relation between Origin and DriveTrain.")
```

Chi2: 40.17843276388697
p-value: 3.975800851556937e-08
Reject H → Origin and DriveTrain are related.

16 Visualizing

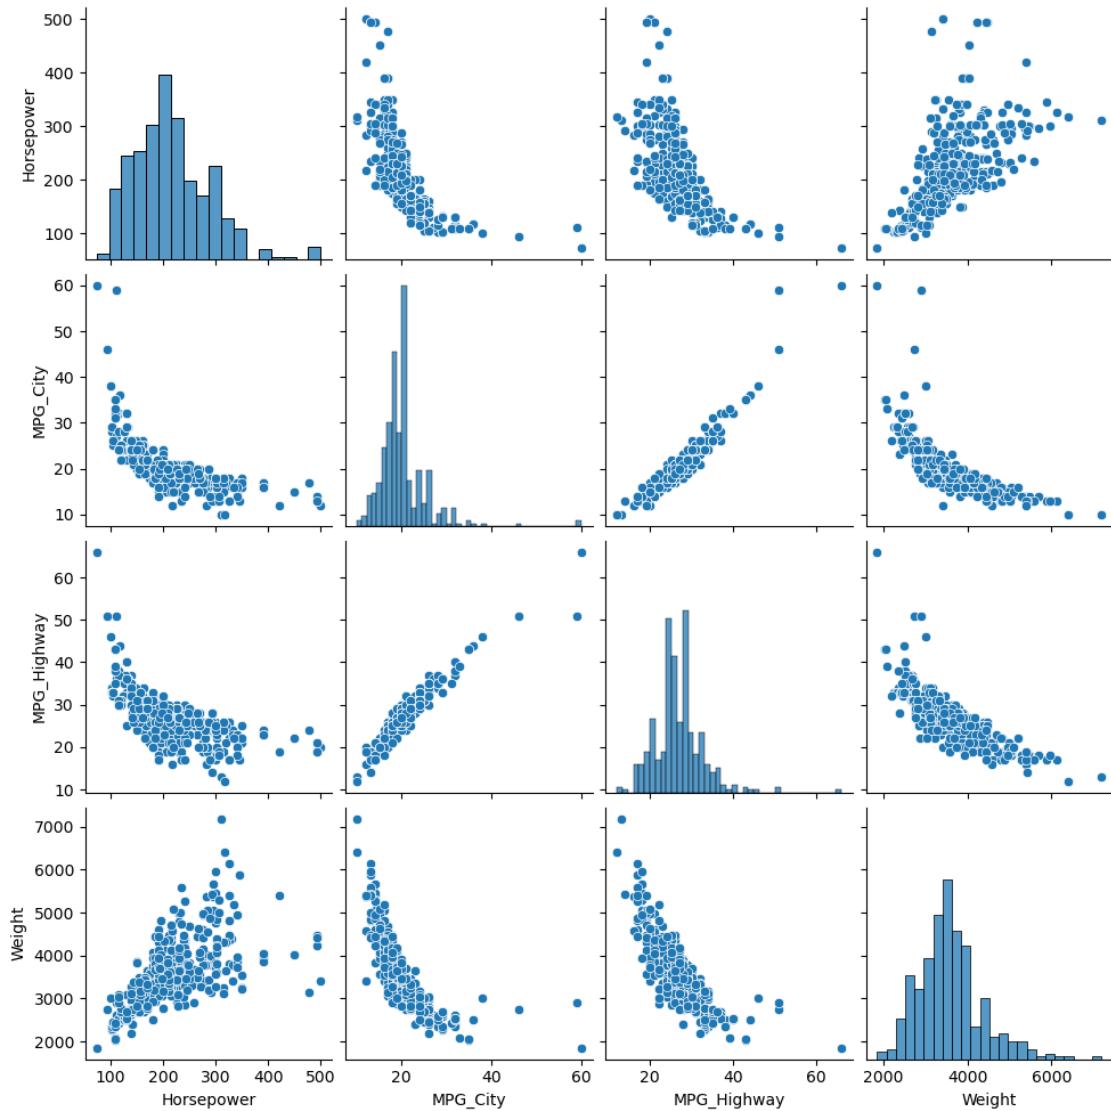
17 Boxplot of MSRP by Origin

```
[35]: sns.boxplot(x='Origin', y='MSRP', data=cars)
plt.title("Car Price (MSRP) by Origin")
plt.show()
```



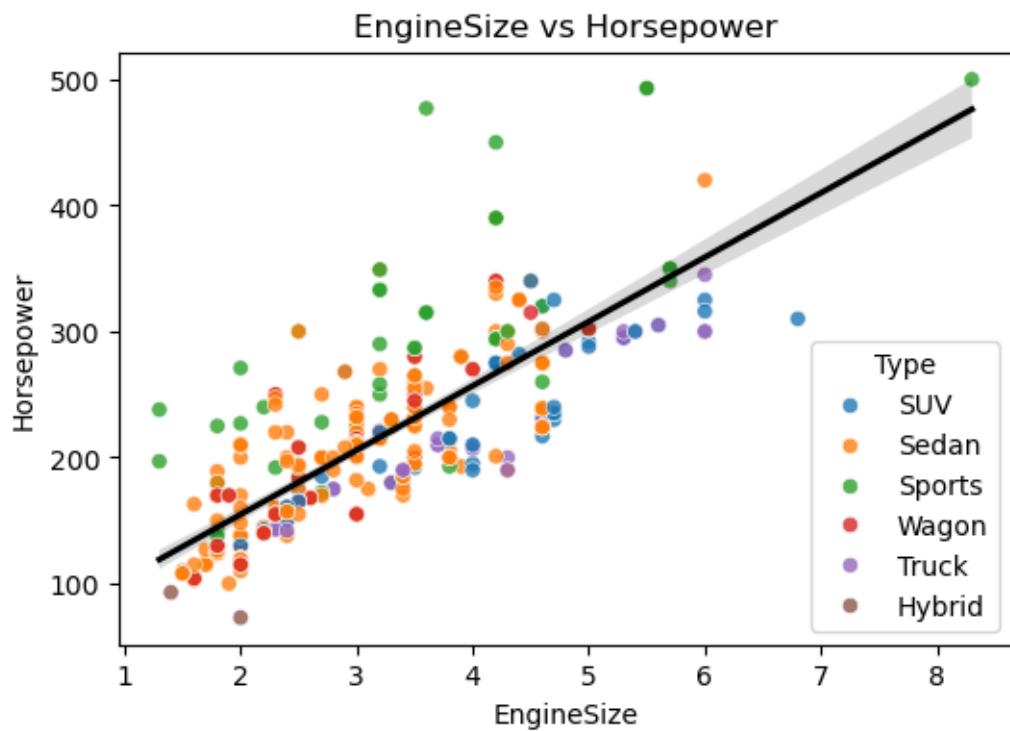
18 Pairplot

```
[36]: sns.pairplot(cars[['Horsepower', 'MPG_City', 'MPG_Highway', 'Weight']])
plt.show()
```



19 scatterplot

```
[39]: plt.figure(figsize=(6,4))
sns.scatterplot(x='EngineSize', y='Horsepower', hue='Type', data=cars, alpha=0.8)
sns.regplot(x='EngineSize', y='Horsepower', data=cars, scatter=False, color='black')
plt.title("EngineSize vs Horsepower")
plt.show()
```



[]: